

# Liang Wang

☎ (+86)157-7413-1055 | ✉ [iggiewang@gmail.com](mailto:iggiewang@gmail.com) | 🌐 [iggiewang.cn](http://iggiewang.cn) | 🌐 [hey-kong](https://hey-kong.com)

## Education

**Huazhong University of Science and Technology**

MSc in Computer Architecture

September 2021 – June 2024 (Expected)

Wuhan, China

**Wuhan University**

B. Eng in Software Engineering, GPA: 3.83/4.0, Rank: 16/258

September 2017 – June 2021

Wuhan, China

## Papers

**Shoggoth: Towards Efficient Edge-Cloud Collaborative Real-Time Video Inference via Adaptive Online Learning**

- **Liang Wang\***, Kai Lu\*, Nan Zhang, Xiaoyang Qu, Jianzong Wang, Jiguang Wan, Guokuan Li, Jing Xiao
- Submitted to *DAC 2023*, accepted

## Internship Experience

**Ping An Technology (Shenzhen) Co., Ltd.**

Algorithm Engineer Intern

February 2022 – Present

Shenzhen, China

- Constructed the **KubeEdge** cluster; developed the custom Mapper to parse and format messages to facilitate communication between IoT devices and the edge; interfaced the MQTT Broker to TDengine via Kuiper to enable data persistence.
- Deployed **Sedna** on the KubeEdge cluster; created multiple video object detection applications with incremental learning utilizing Sedna edge-cloud collaborative inference and training frameworks to address the issue of data drift.
- Implemented the Unsupervised Hard Example Mining algorithm to automatically obtain examples during video object detection as part of the dataset for incremental learning.
- Developed vision AI applications using the **DeepStream SDK** and deployed using Docker, with each able to process 10–30 video streams in real time when deployed on NVIDIA Jetson boards.

**Huawei Technologies Co., Ltd.**

Cloud Infrastructure Software Engineer Intern

November 2020 – April 2021

Shenzhen, China

- Contributed actively to developing the distributed file system **SFS Turbo 2.0**; implemented the POSIX-compliant lookup and read I/O interfaces.
- Achieved non-blocking Inode ID generation during concurrent file creation, with reference to the distributed ID generator Leaf.
- Improved the system throughput by 2–3x by utilizing memory pool and coroutine libraries and upgrading SPDK/DPDK.
- Generated unit test cases (gtest) and analyzed performance issues (vdbench, mdtest, FlameGraph) for the system.

**PingCAP, Inc.**

PingCAP Talent Plan Training Program

July 2019 – August 2019

Beijing, China

- Designed the index recommendation algorithm, inspired by DB2 Advisor; implemented **Index Advisor** for **TiDB** which supported TiDB client-side interaction.
- Proved the effectiveness of Index Advisor by conducting TPC-DS tests, and demonstrated a 20x performance boost for some SQL queries after creating the recommended index.

## Projects

🔗 **LanguorDB**: LSM-tree Based storage engine from ground-up, inspired by LevelDB. Designed the coarse-grained compaction strategy to reduce write amplification; implemented parallel lookup mechanism and row cache to optimize read efficiency.

🔗 **MayflyCache**: Lightweight implementation of a distributed cache. Used consistent hashing to select nodes for load balancing; implemented HTTP-based communication and Protobuf serialization data transfer; added singleflight to prevent hotspot invalidation.

🔗 **HBtree**: Hybrid DRAM-NVM index, whose structure is learned index ALEX (SIGMOD '20) in DRAM and B+Tree in NVM. Designed hot/cold identification to classify ALEX leaf nodes and construct B+trees from cold nodes to transfer cold data to NVM.

🔗 **House Price Analysis**: Web application based on Vue.js, Flask, and MySQL. Developed the back-end and leveraged Scrapy framework to crawl house prices for dozens of cities; designed a database that satisfied the third normal form.

## Skills

**Programming Languages**: Go, C/C++, Rust, Python

**Tech Skills**: Edge Computing, Database Systems, Key-Value Storage Systems, Distributed Systems

## Miscellaneous

Finished **MIT 6.824** and **CSAPP** computer course labs

Submitted 20+ PRs to well-known open-source software (TiDB, Sedna, etc.)